

Architecting Trustworthy Agents: A Multi-Layered Defense from Alignment to Isolation

Kaifa, Jackal, ZHAO

CSIG, Tencent

Shenzhen, China






Content

- I. Background and Motivation**
- II. Alignment & Constitutional Guardrails**
- III. Proactive Security Auditing**
- IV. Privacy Guardian**
- V. Trustworthy Agent Environment**
- VI. Evolution Loop**
- VII. Conclusion**



I. Background & Motivation

Threats to Autonomous AI Agents




[1] Tencent Cloud Hackathon. <https://zc.tencent.com/>. Oct.2025~Apr. 2026.

[2] Li, Songze, et al. "Odysseus: Jailbreaking Commercial Multimodal LLM-integrated Systems via Dual Steganography." NDSS. 2026.

I. Background & Motivation

Threats to Autonomous AI Agents



II. Alignment & Constitutional Guardrails

Multi-Layered Alignment: From Policy to Behavioral Guardrails

Layer III : Alignment through Self-reflection.


- Self-reflection loop
 - ❑ a compliance checklist

Layer II : Instruction-side Semantic Guardrails.

- Semantic interceptor
 - ❑ Jailbreak
 - ❑ Poisoning
 - ❑ ...


Layer I : The Pre-defined Agent Regulation. ^[3]

- Examples of regulations:
 - ❑ "Vulnerability scanning is permitted; data destruction is strictly prohibited. "
 - ❑ "No harmful context permitted."
 - ❑ ...
- Summarized from rules, redlines, ethical constraints.



III. Proactive Security Auditing



Sentinel Architecture: Enforcing Security via Separation of Duties



IV. Privacy Guardian


DCG Engine: The Synergy of Rules and AI

Data Classification and Grading




IV. Privacy Guardian

Inbound Privacy Shield: Preventing Sensitive




IV. Privacy Guardian


Outbound Defense: Mitigating Exfiltration and Tool Misuse



V. Trustworthy Agent Environment




V. Trustworthy Agent Environment




V. Trustworthy Agent Env

Verifiable Trust: Private Cloud Compute (PCC) Architecture




V. Trustworthy Agent Env

Infrastructure Isolation: Security Sandbox



VI. Evolution Loop

Observability & Continuous Alignment





VI. Evolution Loop

Runtime Observability

Full-stack monitoring platform:

- AI Consumption:
 - # Requests
 - # Tokens (Total & Input & Output \$ Modules)
- Fine-Grained Tracking:
 - # Tool calls
 - Module Status
 - Durations



VI. Evolution Loop

Trajectory-Driven Evolution

Agent Annotation Platform

Feedback

Agent Refinement

Editing intelligent agents

- ✓ Basic Information
Fill in the basic information of the intelligent agent
- ✓ Select template
Select configuration template
- 3 Configuration parameters
Enter configuration parameters

Template: Plan-React-Analysis General Template

* Planning module ?

Plan_ Attack plan

Investigation plan for host detecting network attack (

Module ID: module-df

Using template: Plan node

Configuration Preview

* Execution module ?

Alarm information extraction test React

The content of the original alarm information is extracted and passed to other modules for subsequent analysis.


Module ID: module-E

Use template: ReAct node

Configuration Preview

```
{
  "enable_replan": false,
  "tool_deduplication": true,
  "ignore_step_content": true,
  "output_schema": "{\n  \"$schema\": \"http://...\",
  \"system_prompt\": \"You are the alarm information extraction expert, responsible for accurately extracting information from the following alert message and output it in a structured format.\",
  \"user_prompt\": \"# Task\\nExtract structured information from the following alert message and output it in a structured format.\",
  \"output_prompt\": \"\",
  \"model_id\": \"model-...\",
  \"tool_config\": {},
  \"output_tag\": \"event info\"
}
```

Annotation Comparison Sys



AI Annotation Sys

task-015 Manual annotation

Progress marking: 0/0 Return to list

Task Name: Manual Annotation Test Task Agent: agent-dev-001 Configuration evaluation: Customer service quality evaluation indicators

Trace information

Start time: 2025-12-23 21:31:58

End time: 2025-12-23 21:32:10

Execution time: 12.00s Token consumption: 2,800

state: To be labeled Added to dataset

Trace content

User input (task_content):

```
{
  "question": "Hello, how can I help you?"
}
```

Agent response (result_content):

Standard delivery typically takes 3-5 business days, while expedited delivery takes 1-2 business days.

Execution trace (trace_content):

```
{
  "steps": [
    {
      "content": {
        "role": "user",
        "Message": "Please check today's weather forecast and give me some clothing suggestions."
      },
      "step_id": 1,
      "step_type": "user_input",
      "timestamp": 1766496718000
    },
    {
      "content": {
        "plan": [
          1. Call the weather query API to get the current weather.
          2. Analyze weather data (temperature, humidity, probability of precipitation, etc.)
          3. Provide clothing suggestions based on the weather conditions.
        ],
        "thought": "Users need to check the weather and get clothing suggestions. I need"
      }
    }
  ]
}
```

Label Function Area

Already in the dataset Abandon

Indicator Labeling

Based on the trace content, the following metrics were evaluated.

1 Correctness of answer

Was the customer service response accurate and correct?

type: boolean

请输入 boolean 类型的值

评估理由 (可选)

2 Response Integrity

Does the answer contain complete information?

type: boolean

请输入 boolean 类型的值

评估理由 (可选)

3 Service attitude

Was the tone friendly and professional? (1-5 points)

type: int

请输入 int 类型的值

评估理由 (可选)

Human
Annotation Sys


Tracing

Conclusion

Discussion && Future work

From ReAct to be S.A.F.E.R

- User Input → Plan → Reasoning → Action → Observation → Result
- Standardization → Audit → Fortification → Evaluation → Refinement



Future work

- From *probabilistic safety* to **mathematical safety**
- From *experimental trustworthy* to **verifiable trustworthy**
- Make **agent security** to **secure agent**

Thanks

Q & A

Let's begin our discussion ...



Audio



Idea

